

MobileNet_project

September 29, 2025

1 Melanoma Detection using Transfer Learning with MobileNetV2

1.1 Introduction

Transfer learning has become a powerful approach in computer vision, especially for medical imaging tasks where annotated datasets are often limited. Instead of training models from scratch, pretrained convolutional neural networks (CNNs) can be fine-tuned to leverage knowledge gained from large-scale datasets like ImageNet. This strategy enables faster convergence, reduces computational costs, and often yields higher accuracy in specialized domains such as disease classification.

In this project, transfer learning is applied using the MobileNetV2 architecture to the task of melanoma detection, where early and reliable diagnosis is crucial for improving patient outcomes.

1.2 Data

The dataset used in this project is sourced from [Kaggle's Melanoma Cancer Dataset](#). It contains **13,900 high-quality dermoscopic images**, uniformly resized to **224×224 pixels**, representing both benign and malignant skin lesions.

The dataset was curated to support research in dermatology and computer-aided diagnostics, with the goal of enabling the development of machine learning models that can distinguish between healthy and cancerous tissue. Its diversity of lesion appearances provides a realistic challenge for classification tasks, making it a suitable benchmark for evaluating the effectiveness of transfer learning approaches in medical image analysis.

1.2.1 Data loader

Load data from <https://www.kaggle.com/datasets/bhaveshmittal/melanoma-cancer-dataset/data> for binary task, where

- 0 → 'Benign'
- 1 → 'Malignant'

Melanoma is class 1 (Malignant)

Folder structure (I put everything in 'data' folder):

```
data/  
  train/
```

Benign/

Malignant/

test/

Benign/

Malignant/

```
[60]: # libraries
import tensorflow as tf
import os
import random
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'
os.environ['CUDA_VISIBLE_DEVICES'] = ''

[62]: def load_dataset_simple_streamed(base_dir, img_size=(224, 224), batch_size=16):
    class_names = ['Benign', 'Malignant']
    label_map = {name: idx for idx, name in enumerate(class_names)}

    all_image_paths = []
    all_labels = []

    # Combine all images from train/ and test/ folders
    for subset in ['train', 'test']:
        for class_name in class_names:
            folder = os.path.join(base_dir, subset, class_name)
            label = label_map[class_name]

            for fname in os.listdir(folder):
                if fname.endswith('.jpg'):
                    all_image_paths.append(os.path.join(folder, fname))
                    all_labels.append(label)

    # Shuffle together
    combined = list(zip(all_image_paths, all_labels))
    random.seed(42)
    random.shuffle(combined)
    all_image_paths, all_labels = zip(*combined)

    # Create tf.data.Dataset from paths and labels
    path_ds = tf.data.Dataset.from_tensor_slices(list(all_image_paths))
    label_ds = tf.data.Dataset.from_tensor_slices(list(all_labels))

    def load_and_preprocess(path, label):
        image = tf.io.read_file(path)
        image = tf.image.decode_jpeg(image, channels=3)
```

```

        image = tf.image.resize(image, img_size)
        image = tf.cast(image, tf.float32) / 255.0 # normalization
        return image, label

    full_ds = tf.data.Dataset.zip((path_ds, label_ds))
    full_ds = full_ds.map(load_and_preprocess, num_parallel_calls=tf.data.
↪AUTOTUNE)

    # Split into 70/15/15
    total = len(all_image_paths)
    train_size = int(0.7 * total)
    val_size = int(0.15 * total)

    train_ds = full_ds.take(train_size).batch(batch_size).prefetch(tf.data.
↪AUTOTUNE)
    val_ds = full_ds.skip(train_size).take(val_size).batch(batch_size).
↪prefetch(tf.data.AUTOTUNE)
    test_ds = full_ds.skip(train_size + val_size).batch(batch_size).prefetch(tf.
↪data.AUTOTUNE)

    return train_ds, val_ds, test_ds, class_names

```

```
[27]: train_ds, val_ds, test_ds, class_names = load_dataset_simple_streamed("data")
```

```
[29]: for images, labels in train_ds.take(1):
        print("Train batch shape:", images.shape, labels.shape)

    for images, labels in val_ds.take(1):
        print("Val batch shape:", images.shape, labels.shape)

    for images, labels in test_ds.take(1):
        print("Test batch shape:", images.shape, labels.shape)

```

Train batch shape: (16, 224, 224, 3) (16,)

Val batch shape: (16, 224, 224, 3) (16,)

Test batch shape: (16, 224, 224, 3) (16,)

```
[31]: def count_samples(ds):
        return sum(1 for _ in ds.unbatch())

    print("Dataset sizes:")
    print("Train size:", count_samples(train_ds))
    print("Val size: ", count_samples(val_ds))
    print("Test size: ", count_samples(test_ds))

```

Dataset sizes:

Train size: 9715


Val size: 2081

Test size: 2083

```
[34]: import matplotlib.pyplot as plt

def plot_images_from_ds(ds, class_names=None, n=10):
    plt.figure(figsize=(15, 3))
    for i, (image, label) in enumerate(ds.unbatch().take(n)):
        plt.subplot(1, n, i + 1)
        plt.imshow(image.numpy())
        label_val = label.numpy()
        plt.title(class_names[label_val] if class_names else str(label_val),
        ↪fontsize=10)
        plt.axis('off')
    plt.tight_layout()
    plt.show()
```

```
[47]: plot_images_from_ds(train_ds, class_names)
```



```
[49]: # Check image size and dataset size
from PIL import Image
import os

for phase in ['train', 'test']:
    for label in ['Benign', 'Malignant']:
        path = f"./data/{phase}/{label}"
        sample_image = os.listdir(path)[0]
        with Image.open(os.path.join(path, sample_image)) as img:
            print(f"{phase}/{label} - Sample image size: {img.size}, Total
            ↪images: {len(os.listdir(path))}")
```

```
train/Benign - Sample image size: (224, 224), Total images: 6289
train/Malignant - Sample image size: (224, 224), Total images: 5590
test/Benign - Sample image size: (224, 224), Total images: 1000
test/Malignant - Sample image size: (224, 224), Total images: 1000
```

```
[53]: import numpy as np

def check_class_distribution_ds(train_ds, val_ds, test_ds, class_names):
    def get_counts(ds):
        labels = []
        for _, y in ds.unbatch():
```

```

        labels.append(int(y.numpy()))
    unique, counts = np.unique(labels, return_counts=True)
    return dict(zip(unique, counts)), len(labels)

    for name, ds in zip(['Train', 'Validation', 'Test'], [train_ds, val_ds,
↪test_ds]):
        counts, total = get_counts(ds)
        print(f"\n{name} set:")
        for i, cls in enumerate(class_names):
            count = counts.get(i, 0)
            print(f"  {cls}: {count} ({100 * count / total:.2f}%)")

```

```
[56]: check_class_distribution_ds(train_ds, val_ds, test_ds, class_names)
```

Train set:

Benign: 5087 (52.36%)
Malignant: 4628 (47.64%)

Validation set:

Benign: 1104 (53.05%)
Malignant: 977 (46.95%)

Test set:

Benign: 1098 (52.71%)
Malignant: 985 (47.29%)

1.3 Model Architecture - MobileNetV2

MobileNetV2 is a lightweight and efficient convolutional neural network architecture proposed by Sandler et al. (2018), optimized for mobile and embedded vision applications. It introduces two key innovations: inverted residual blocks and linear bottlenecks, which significantly reduce computational cost and model size while preserving accuracy. Instead of traditional residual connections, MobileNetV2 expands the number of channels, applies depthwise separable convolutions, and then projects back to a lower-dimensional space—hence the term “inverted.” This architecture maintains high representational power with fewer parameters, making it ideal for scenarios where computational efficiency is critical, such as mobile and edge devices.

In this project, I implemented and evaluated three configurations of the MobileNetV2 convolutional neural network, all based on a common architecture that utilized the pretrained MobileNetV2 backbone as a feature extractor. The base model was initialized with ImageNet weights and configured without the top classification layer (`include_top=False`). A custom classification head was appended to adapt the model for binary classification. All models were compiled using the Adam optimizer with a learning rate of 0.001, the binary cross-entropy loss function (suitable for binary classification tasks), and accuracy as the primary evaluation metric.

Architecture details:

- **Input Preprocessing:** Images resized to 224×224×3 and normalized to [0,1].
- **Backbone:** Pretrained MobileNetV2 (`weights='imagenet', include_top=False`).

- Custom Classification Head:
- GlobalAveragePooling2D
- Dense(128, activation='relu')
- Dropout(0.3)
- Dense(1, activation='sigmoid')

Training setup:

- Optimizer: Adam (learning rate = 0.001)
- Loss function: Binary Cross-Entropy
- Metric: Accuracy

1.4 Experiments

Three experimental configurations were tested to evaluate the effect of transfer learning and fine-tuning:

1. Fully frozen base model
2. Partially fine-tuned base model (last 30 layers unfrozen)
3. Fine-tuned base with data augmentation

1.4.1 1. Fully frozen base model with custom classification head

```
[74]: from tensorflow.keras.applications import MobileNetV2
      from tensorflow.keras import layers, models

      # Load pretrained MobileNetV2 base
      base_model = MobileNetV2(input_shape=(224, 224, 3), include_top=False,
                               weights='imagenet')
      base_model.trainable = False # Freeze all layers

      # Add custom classification head
      model = models.Sequential([
          base_model,
          layers.GlobalAveragePooling2D(),
          layers.Dense(128, activation='relu'),
          layers.Dropout(0.3),
          layers.Dense(1, activation='sigmoid') # Binary classification
      ])
```

```
[76]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2,257,984

global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 1280)	0
dense_2 (Dense)	(None, 128)	163,968
dropout_1 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 1)	129

Total params: 2,422,081 (9.24 MB)

Trainable params: 164,097 (641.00 KB)

Non-trainable params: 2,257,984 (8.61 MB)

```
[96]: model.compile(
        optimizer='adam',
        loss='binary_crossentropy',
        metrics=['accuracy']
    )
```


```
[108]: import matplotlib.pyplot as plt

# Train the model
start_train = time.time()
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=10
)
end_train = time.time()
train_duration = end_train - start_train
print(f"\n Training time: {train_duration:.2f} seconds")

# Plot training & validation loss
plt.figure(figsize=(8, 5))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.grid(True)
plt.show()
```

Epoch 1/10
608/608 67s 110ms/step -
accuracy: 0.8985 - loss: 0.2495 - val_accuracy: 0.8779 - val_loss: 0.2953
Epoch 2/10
608/608 69s 114ms/step -
accuracy: 0.9089 - loss: 0.2335 - val_accuracy: 0.8698 - val_loss: 0.2992
Epoch 3/10
608/608 65s 107ms/step -
accuracy: 0.9112 - loss: 0.2248 - val_accuracy: 0.8693 - val_loss: 0.3033
Epoch 4/10
608/608 65s 108ms/step -
accuracy: 0.9154 - loss: 0.2133 - val_accuracy: 0.8698 - val_loss: 0.3298
Epoch 5/10
608/608 67s 111ms/step -
accuracy: 0.9193 - loss: 0.2015 - val_accuracy: 0.8784 - val_loss: 0.3031
Epoch 6/10
608/608 67s 111ms/step -
accuracy: 0.9240 - loss: 0.1961 - val_accuracy: 0.8765 - val_loss: 0.3146
Epoch 7/10
608/608 67s 110ms/step -
accuracy: 0.9296 - loss: 0.1792 - val_accuracy: 0.8784 - val_loss: 0.3193
Epoch 8/10
608/608 66s 108ms/step -
accuracy: 0.9310 - loss: 0.1720 - val_accuracy: 0.8779 - val_loss: 0.3256
Epoch 9/10
608/608 66s 108ms/step -
accuracy: 0.9396 - loss: 0.1600 - val_accuracy: 0.8746 - val_loss: 0.3364
Epoch 10/10
608/608 66s 109ms/step -
accuracy: 0.9442 - loss: 0.1520 - val_accuracy: 0.8722 - val_loss: 0.3516

Training time: 666.30 seconds



Evaluation

```
[121]: import numpy as np
import matplotlib.pyplot as plt
import time
from sklearn.metrics import (
    confusion_matrix, ConfusionMatrixDisplay, classification_report,
    accuracy_score, f1_score, precision_score, recall_score
)

def evaluate_classification_model(model, train_ds, test_ds, class_names=None):
    def extract_labels(ds):
        y_true = []
        for _, label in ds.unbatch():
            y_true.append(int(label.numpy()))
        return np.array(y_true)

    def get_predictions(ds, label=""):
        start_predict = time.time()
        y_pred_probs = model.predict(ds, verbose=0).flatten()
        end_predict = time.time()

        predict_duration = end_predict - start_predict
```

```

    print(f"\nPrediction time on {label} set: {predict_duration:.2f}\n↪seconds")
    y_pred = (y_pred_probs > 0.5).astype(np.int64)
    return y_pred

def evaluate(y_true, y_pred, label=""):
    acc = accuracy_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred)
    rec = recall_score(y_true, y_pred)

    print(f"\n{n{label} SET METRICS")
    print(f"  Accuracy:      {acc:.4f}")
    print(f"  F1 Score:       {f1:.4f}")
    print(f"  Precision:      {prec:.4f}")
    print(f"  Recall:         {rec:.4f}")
    print(f"\nClassification Report ({label}):")
    print(classification_report(
        y_true, y_pred,
        target_names=class_names if class_names else None,
        zero_division=0
    ))

    return acc, f1, y_true, y_pred

# Ground truth
y_true_train = extract_labels(train_ds)
y_true_test = extract_labels(test_ds)

# Predictions
y_pred_train = get_predictions(train_ds, "TRAIN" )
y_pred_test = get_predictions(test_ds, "TEST")

# Evaluate
train_acc, train_f1, y_train, y_train_pred = evaluate(y_true_train, ↪
↪y_pred_train, "TRAIN")
test_acc, test_f1, y_test, y_test_pred = evaluate(y_true_test, y_pred_test, ↪
↪"TEST")

# Plot confusion matrices
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

ConfusionMatrixDisplay.from_predictions(
    y_train, y_train_pred,
    normalize='true', ax=ax1, display_labels=class_names
)

```

```

ax1.set_title(f'Train Confusion Matrix\nAcc: {train_acc:.2f}, F1: {train_f1:
↪.2f}')
ax1.tick_params(axis='x', rotation=45)

ConfusionMatrixDisplay.from_predictions(
    y_test, y_test_pred,
    normalize='true', ax=ax2, display_labels=class_names
)
ax2.set_title(f'Test Confusion Matrix\nAcc: {test_acc:.2f}, F1: {test_f1:.
↪.2f}')
ax2.tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()

```

```

[123]: evaluate_classification_model(model, train_ds, test_ds, class_names=['Benign',
↪ 'Malignant'])

```

Prediction time on TRAIN set: 52.40 seconds

Prediction time on TEST set: 12.16 seconds

TRAIN SET METRICS

```

Accuracy:      0.9425
F1 Score:      0.9397
Precision:     0.9386
Recall:        0.9408

```

Classification Report (TRAIN):

	precision	recall	f1-score	support
Benign	0.95	0.94	0.94	5087
Malignant	0.94	0.94	0.94	4628
accuracy			0.94	9715
macro avg	0.94	0.94	0.94	9715
weighted avg	0.94	0.94	0.94	9715

TEST SET METRICS

```

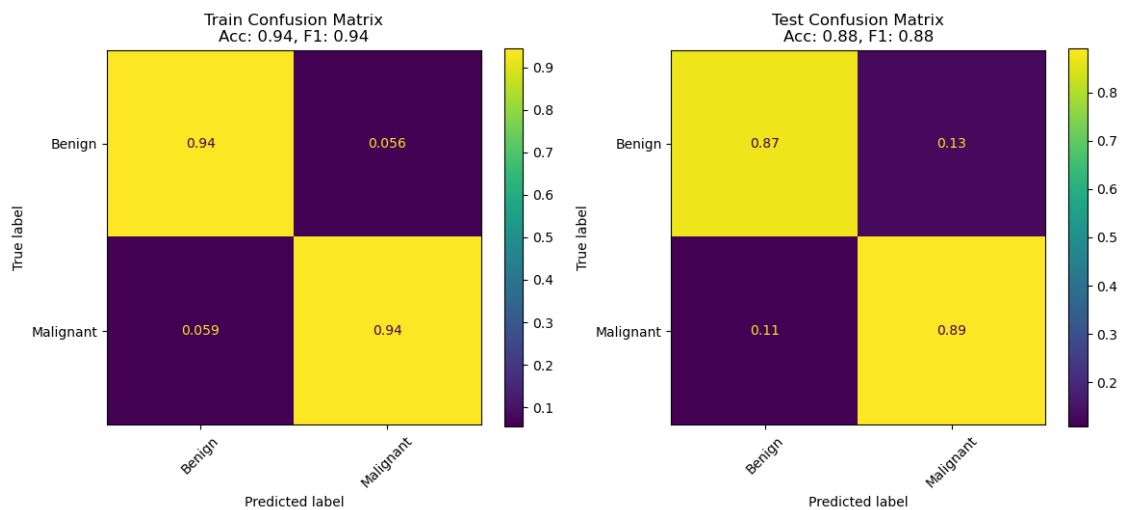
Accuracy:      0.8805
F1 Score:      0.8757
Precision:     0.8615
Recall:        0.8904

```

Classification Report (TEST):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Benign	0.90	0.87	0.88	1098
Malignant	0.86	0.89	0.88	985
accuracy			0.88	2083
macro avg	0.88	0.88	0.88	2083
weighted avg	0.88	0.88	0.88	2083



```
[129]: model.save_weights("fine_tuned_final.weights.h5")
```

1.4.2 2. Partially fine-tuned base model (last 30 layers unfrozen)

```
[131]: # Unfreeze from a specific layer
base_model.trainable = True
for layer in base_model.layers[:-30]: # Freeze all except last 30
    layer.trainable = False

# Recompile
model.compile(
    optimizer=tf.keras.optimizers.Adam(1e-5), #smaller learning rate
    loss='binary_crossentropy',
    metrics=['accuracy']
)

# Fine-tune
start_train = time.time()
history_finetune = model.fit(
    train_ds,
    validation_data=val_ds,
```

```

        epochs=10
    )
    end_train = time.time()
    train_duration = end_train - start_train
    print(f"\n Training time: {train_duration:.2f} seconds")

    # Plot training & validation loss
    plt.figure(figsize=(8, 5))
    plt.plot(history_finetune.history['loss'], label='Training Loss')
    plt.plot(history_finetune.history['val_loss'], label='Validation Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.title('Training and Validation Loss')
    plt.legend()
    plt.grid(True)
    plt.show()


```

```

Epoch 1/10
608/608          93s 148ms/step -
accuracy: 0.7542 - loss: 1.2247 - val_accuracy: 0.8496 - val_loss: 0.6041
Epoch 2/10
608/608          89s 146ms/step -
accuracy: 0.8563 - loss: 0.3947 - val_accuracy: 0.8611 - val_loss: 0.4549
Epoch 3/10
608/608          90s 148ms/step -
accuracy: 0.9014 - loss: 0.2504 - val_accuracy: 0.8693 - val_loss: 0.3964
Epoch 4/10
608/608          90s 148ms/step -
accuracy: 0.9178 - loss: 0.2014 - val_accuracy: 0.8731 - val_loss: 0.3688
Epoch 5/10
608/608          89s 147ms/step -
accuracy: 0.9369 - loss: 0.1517 - val_accuracy: 0.8751 - val_loss: 0.3557
Epoch 6/10
608/608          89s 147ms/step -
accuracy: 0.9574 - loss: 0.1113 - val_accuracy: 0.8746 - val_loss: 0.3651
Epoch 7/10
608/608          88s 145ms/step -
accuracy: 0.9692 - loss: 0.0830 - val_accuracy: 0.8779 - val_loss: 0.3674
Epoch 8/10
608/608          88s 146ms/step -
accuracy: 0.9838 - loss: 0.0549 - val_accuracy: 0.8803 - val_loss: 0.4086
Epoch 9/10
608/608          90s 148ms/step -
accuracy: 0.9871 - loss: 0.0395 - val_accuracy: 0.8789 - val_loss: 0.4109
Epoch 10/10
608/608          92s 151ms/step -
accuracy: 0.9926 - loss: 0.0282 - val_accuracy: 0.8770 - val_loss: 0.4598

```

Training time: 898.50 seconds



Evaluation

```
[133]: evaluate_classification_model(model, train_ds, test_ds, class_names=['Benign', 'Malignant'])
```

```
2025-05-28 16:27:01.920955: I tensorflow/core/framework/local_rendezvous.cc:407] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence
```

Prediction time on TRAIN set: 53.78 seconds

Prediction time on TEST set: 12.14 seconds

TRAIN SET METRICS

Accuracy:	0.9686
F1 Score:	0.9671
Precision:	0.9653
Recall:	0.9689

Classification Report (TRAIN):

	precision	recall	f1-score	support
Benign	0.97	0.97	0.97	5087

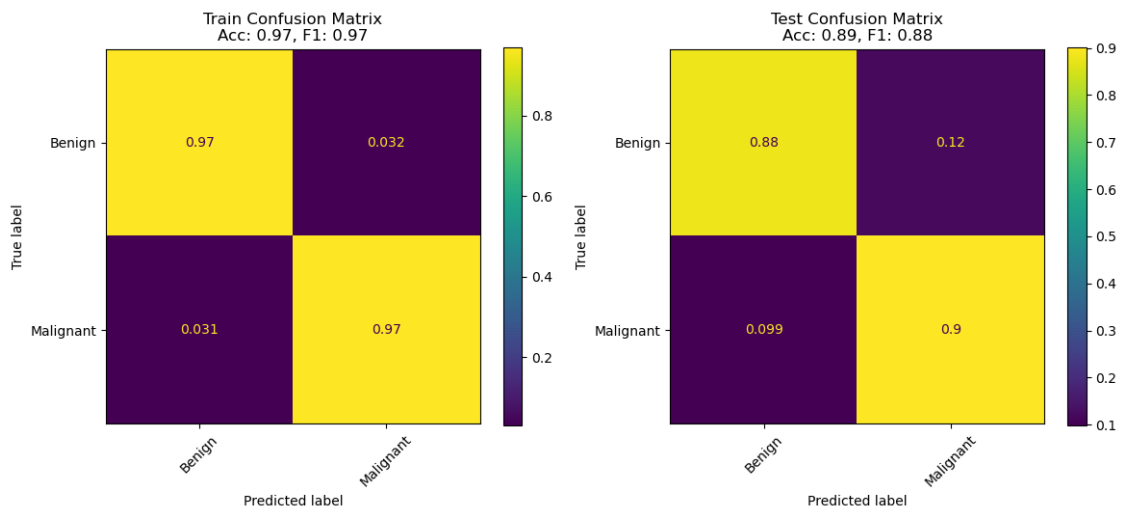
Malignant	0.97	0.97	0.97	4628
accuracy			0.97	9715
macro avg	0.97	0.97	0.97	9715
weighted avg	0.97	0.97	0.97	9715

TEST SET METRICS

Accuracy: 0.8891
 F1 Score: 0.8848
 Precision: 0.8696
 Recall: 0.9005

Classification Report (TEST):

	precision	recall	f1-score	support
Benign	0.91	0.88	0.89	1098
Malignant	0.87	0.90	0.88	985
accuracy			0.89	2083
macro avg	0.89	0.89	0.89	2083
weighted avg	0.89	0.89	0.89	2083



```
[135]: model.save_weights("unfreezed_final.weights.h5")
```

1.4.3 3. Fine-tuned base with data augmentation

```
[137]: from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau

early_stop = EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)

reduce_lr = ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.5,
    patience=5,
    min_lr=1e-6
)
```

Augmentation

```
[167]: import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras import layers

# Define augmentation
data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal_and_vertical"), # Randomly flips images
    ↪horizontally and vertically
    layers.RandomRotation(0.5), # Rotates images randomly by up to 50% of 360°
    layers.RandomZoom(0.2), # Zooms in or out randomly by up to 20%
    layers.RandomTranslation(0.1, 0.1), #Shifts image horizontally and
    ↪vertically by 10%
])

# Get a sample image from train_ds
for images, labels in train_ds.take(1):
    sample_image = images[0]
    break

# Plot 8 augmented + 1 original
plt.figure(figsize=(10, 10))
for i in range(9):
    ax = plt.subplot(3, 3, i + 1)

    if i < 8:
        # Force training=True so augmentations apply
        augmented_img = data_augmentation(tf.expand_dims(sample_image, 0),
        ↪training=True)[0]
        img = tf.clip_by_value(augmented_img, 0.0, 1.0)
```



```


        title = f"Augmented {i + 1}"
    else:
        img = sample_image
        title = "Original"

    plt.imshow(img.numpy())
    plt.title(title)
    plt.axis("off")

plt.tight_layout()
plt.show()

```


Augmented 1




Augmented 2




Augmented 3




Augmented 4




Augmented 5




Augmented 6




Augmented 7



Augmented 8



Original



Adding augmentation layer

```
[169]: from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras import layers, models

# Load pretrained MobileNetV2 base
base_model = MobileNetV2(input_shape=(224, 224, 3), include_top=False,
    weights='imagenet')
base_model.trainable = False # Freeze all layers

# Add custom classification head
model = models.Sequential([
    data_augmentation,
    base_model,
    layers.GlobalAveragePooling2D(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.3),
    layers.Dense(1, activation='sigmoid') # Binary classification
])

[171]: model.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
sequential_9 (Sequential)	(1, 224, 224, 3)	0
mobilenetv2_1.00_224 (Functional)	(1, 7, 7, 1280)	2,257,984
global_average_pooling2d_2 (GlobalAveragePooling2D)	(1, 1280)	0
dense_4 (Dense)	(1, 128)	163,968
dropout_2 (Dropout)	(1, 128)	0
dense_5 (Dense)	(1, 1)	129

Total params: 2,422,081 (9.24 MB)

Trainable params: 164,097 (641.00 KB)

Non-trainable params: 2,257,984 (8.61 MB)

```
[173]: import matplotlib.pyplot as plt
import time

# Unfreeze from a specific layer
base_model.trainable = True
for layer in base_model.layers[:-30]: # Freeze all except last 30
    layer.trainable = False

# Recompile
model.compile(
    optimizer=tf.keras.optimizers.Adam(1e-5), #smaller learning rate
    loss='binary_crossentropy',
    metrics=['accuracy']
)

batch_size=64
# Fine-tune
start_train = time.time()
history_finetune = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=30,
    callbacks=[early_stop]
)

end_train = time.time()
train_duration = end_train - start_train
print(f"\n Training time: {train_duration:.2f} seconds")

# Plot training & validation loss
plt.figure(figsize=(8, 5))
plt.plot(history_finetune.history['loss'], label='Training Loss')
plt.plot(history_finetune.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.grid(True)
plt.show()
```

Epoch 1/30

608/608 94s 149ms/step -

accuracy: 0.6815 - loss: 0.5764 - val_accuracy: 0.7876 - val_loss: 0.4327

Epoch 2/30

608/608 90s 147ms/step -

accuracy: 0.8049 - loss: 0.4169 - val_accuracy: 0.8275 - val_loss: 0.3702

Epoch 3/30
608/608 94s 155ms/step -
accuracy: 0.8233 - loss: 0.3881 - val_accuracy: 0.8496 - val_loss: 0.3401
Epoch 4/30
608/608 94s 155ms/step -
accuracy: 0.8335 - loss: 0.3733 - val_accuracy: 0.8525 - val_loss: 0.3298
Epoch 5/30
608/608 95s 156ms/step -
accuracy: 0.8501 - loss: 0.3554 - val_accuracy: 0.8544 - val_loss: 0.3298
Epoch 6/30
608/608 96s 158ms/step -
accuracy: 0.8541 - loss: 0.3390 - val_accuracy: 0.8467 - val_loss: 0.3319
Epoch 7/30
608/608 97s 160ms/step -
accuracy: 0.8540 - loss: 0.3341 - val_accuracy: 0.8578 - val_loss: 0.3188
Epoch 8/30
608/608 95s 157ms/step -
accuracy: 0.8567 - loss: 0.3292 - val_accuracy: 0.8587 - val_loss: 0.3135
Epoch 9/30
608/608 93s 153ms/step -
accuracy: 0.8617 - loss: 0.3198 - val_accuracy: 0.8630 - val_loss: 0.3062
Epoch 10/30
608/608 92s 152ms/step -
accuracy: 0.8645 - loss: 0.3162 - val_accuracy: 0.8635 - val_loss: 0.3071
Epoch 11/30
608/608 94s 154ms/step -
accuracy: 0.8734 - loss: 0.3069 - val_accuracy: 0.8597 - val_loss: 0.3122
Epoch 12/30
608/608 93s 153ms/step -
accuracy: 0.8672 - loss: 0.3056 - val_accuracy: 0.8606 - val_loss: 0.3112
Epoch 13/30
608/608 95s 156ms/step -
accuracy: 0.8717 - loss: 0.3020 - val_accuracy: 0.8630 - val_loss: 0.3061
Epoch 14/30
608/608 93s 153ms/step -
accuracy: 0.8749 - loss: 0.2951 - val_accuracy: 0.8717 - val_loss: 0.2946
Epoch 15/30
608/608 94s 155ms/step -
accuracy: 0.8765 - loss: 0.2910 - val_accuracy: 0.8640 - val_loss: 0.3057
Epoch 16/30
608/608 95s 155ms/step -
accuracy: 0.8842 - loss: 0.2777 - val_accuracy: 0.8707 - val_loss: 0.2974
Epoch 17/30
608/608 94s 155ms/step -
accuracy: 0.8840 - loss: 0.2815 - val_accuracy: 0.8717 - val_loss: 0.2927
Epoch 18/30
608/608 94s 155ms/step -
accuracy: 0.8865 - loss: 0.2727 - val_accuracy: 0.8640 - val_loss: 0.3038

Epoch 19/30
608/608 93s 153ms/step -
accuracy: 0.8858 - loss: 0.2731 - val_accuracy: 0.8626 - val_loss: 0.3101
Epoch 20/30
608/608 95s 156ms/step -
accuracy: 0.8871 - loss: 0.2663 - val_accuracy: 0.8779 - val_loss: 0.2887
Epoch 21/30
608/608 92s 151ms/step -
accuracy: 0.8870 - loss: 0.2653 - val_accuracy: 0.8799 - val_loss: 0.2802
Epoch 22/30
608/608 94s 154ms/step -
accuracy: 0.8921 - loss: 0.2590 - val_accuracy: 0.8659 - val_loss: 0.3165
Epoch 23/30
608/608 94s 155ms/step -
accuracy: 0.8935 - loss: 0.2519 - val_accuracy: 0.8784 - val_loss: 0.2953
Epoch 24/30
608/608 94s 155ms/step -
accuracy: 0.8906 - loss: 0.2568 - val_accuracy: 0.8760 - val_loss: 0.2915
Epoch 25/30
608/608 93s 153ms/step -
accuracy: 0.9007 - loss: 0.2462 - val_accuracy: 0.8707 - val_loss: 0.3074
Epoch 26/30
608/608 92s 152ms/step -
accuracy: 0.8924 - loss: 0.2516 - val_accuracy: 0.8794 - val_loss: 0.2860
Epoch 27/30
608/608 92s 151ms/step -
accuracy: 0.9007 - loss: 0.2465 - val_accuracy: 0.8775 - val_loss: 0.2917
Epoch 28/30
608/608 91s 150ms/step -
accuracy: 0.8998 - loss: 0.2363 - val_accuracy: 0.8760 - val_loss: 0.3003
Epoch 29/30
608/608 93s 153ms/step -
accuracy: 0.9075 - loss: 0.2290 - val_accuracy: 0.8827 - val_loss: 0.2803
Epoch 30/30
608/608 94s 155ms/step -
accuracy: 0.9030 - loss: 0.2386 - val_accuracy: 0.8799 - val_loss: 0.2907

Training time: 2809.18 seconds



Evaluation

```
[182]: evaluate_classification_model(model, train_ds, test_ds, class_names=['Benign', 'Malignant'])
```

Prediction time on TRAIN set: 50.10 seconds

Prediction time on TEST set: 11.14 seconds

TRAIN SET METRICS

```
Accuracy:      0.8938
F1 Score:      0.8828
Precision:     0.9301
Recall:        0.8401
```

Classification Report (TRAIN):

	precision	recall	f1-score	support
Benign	0.87	0.94	0.90	5087
Malignant	0.93	0.84	0.88	4628
accuracy			0.89	9715
macro avg	0.90	0.89	0.89	9715

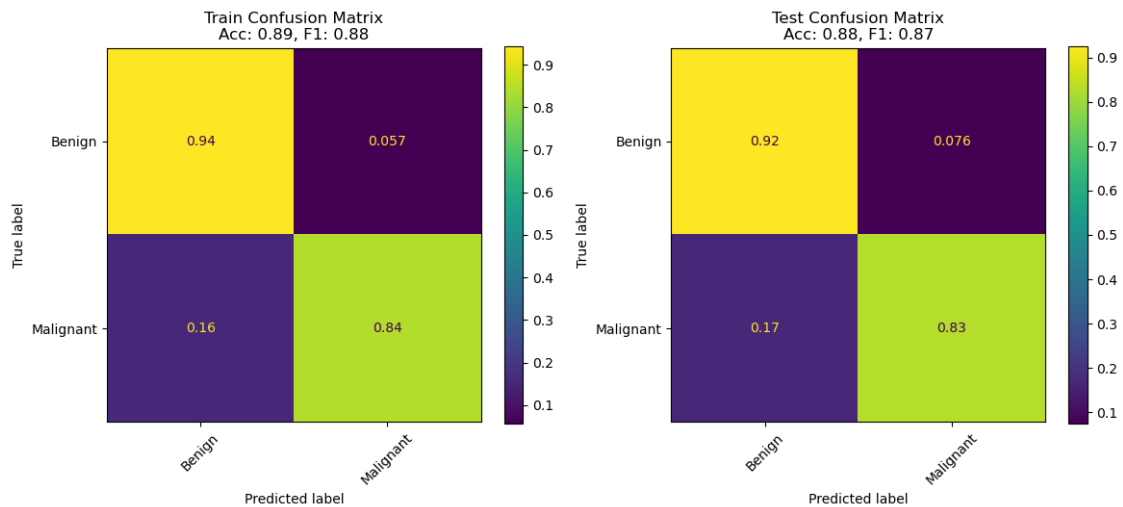
weighted avg	0.90	0.89	0.89	9715
--------------	------	------	------	------

TEST SET METRICS

Accuracy:	0.8805
F1 Score:	0.8680
Precision:	0.9080
Recall:	0.8315

Classification Report (TEST):

	precision	recall	f1-score	support
Benign	0.86	0.92	0.89	1098
Malignant	0.91	0.83	0.87	985
accuracy			0.88	2083
macro avg	0.88	0.88	0.88	2083
weighted avg	0.88	0.88	0.88	2083



```
[184]: model.save_weights("unfreezed_augmented_final.weights.h5")
```

1.5 Conclusion

We evaluated three different transfer learning strategies with MobileNetV2 for melanoma classification. The results are summarized below:

Configuration	Epochs	Accuracy	F1-score	Recall (Malignant)
Frozen MobileNetV2 base + custom head	10	88.05%	0.8757	0.89

Configuration	Epochs	Accuracy	F1-score	Recall (Malignant)
Partially fine-tuned MobileNetV2 (last 30 layers unfrozen)	10	88.91%	0.8848	0.90
Fine-tuned MobileNetV2 with data augmentation	30	88.05%	0.8680	0.83

1.5.1 Analysis

- The **frozen base model with only a custom classification head** produced a strong baseline, transferring general ImageNet features but showing limited adaptation to the unique visual characteristics of melanoma lesions.
- The **partially fine-tuned model (last 30 layers unfrozen)** provided the best results. Allowing deeper layers to adapt to lesion-specific features improved generalization and delivered the highest recall, which is especially critical for medical applications where false negatives must be minimized.
- The **fine-tuned model with additional data augmentation**, even when trained for longer, demonstrated steady learning but did not surpass the selectively fine-tuned configuration. The increased complexity and regularization from augmentation likely slowed convergence and reduced sensitivity to malignant cases.

1.5.2 Final Conclusion

Among the three strategies, the **partially fine-tuned MobileNetV2** demonstrated the most effective trade-off between accuracy, F1-score, and recall. This highlights the importance of selective fine-tuning in transfer learning, as it allows the model to adapt deeper layers to domain-specific features without over-regularization. For melanoma detection tasks, where minimizing false negatives is critical, this configuration provides the most reliable results.

1.6 Contributions

This project contributes by:

- Applying and comparing three distinct MobileNetV2 configurations on a curated melanoma dataset.
- Emphasizing **recall** as a critical metric for clinical tasks.
- Analyzing the impact of **layer unfreezing, data augmentation, and training duration**.
- Providing practical insights for deploying lightweight CNNs in sensitive domains like dermatology.